

ENVS 193DS Final

Stephen Bond

2026-06-09

Table of contents

Set Up	2
Problem 1. Research Writing	3
a. Transparent Statistical Methods	3
b. Suggestions for rewriting	3
Part 1.	3
Part 2.	3
c. Further Analysis	3
Problem 2. Data Analysis	4
a. Clean and Wrangle	4
b. Response Variable	5
c. Purpose of Study	5
Tree Height	5
Acacia Ant Species	5
d. Table of Models	6
e. Run the Models	6
f. Select the Best Model	6
g. Check the Diagnostics	7
h. Visualize the Model Predictions	7
i. Write a Caption for your Figure	9
j. Calculate Model Predictions	10
k. Interpret your Results	10
Problem 3. Affective and Exploratory Visualizations	11
a. Comparing Visualizations	11
How are the visualizations different from each other in the way you have represented your data?	11
What similarities do you see between all your visualizations?	11

What patterns (e.g. differences in means/counts/proportions/medians, trends through time, relationships between variables) do you see in each visualization? Are these different between visualizations? If so, why? If not, why not?	12
What kinds of feedback did you get during week 9 workshop or from the instructors? How did you implement or try those suggestions? If you tried and kept those suggestions, explain how and why; if not, explain why not?	12
b. Designing an Analysis	12
c. Check any Assumptions and Run your Analysis	13
Write in Model	13
Check the Diagnostics with DHARMA	13
Model Coefficients and Calculating Probabilities	14
d. Create a Visualization	16
e. Write a Caption	17
f. Write About Your Results	17
g. Sharing Your Affective Visualization	18
GitHub Repo	

Set Up

```
# loading in packages
library(tidyverse) # general use
library(janitor) # cleaning data
library(here) # organizing files
library(rstatix) # statistical relationships
library(MuMIn) # modelling
library(DHARMA) # residual diagnostic simulation
library(ggeffects) # model predictions
library(ggplot2) # visualizations

# read in data nest data
nests <- read_csv(here("data", "nest_data_final.csv"))

# read in and make object for personal data
dinner <- read_csv(here("data", "Dinner_Data.csv"))
```

Problem 1. Research Writing

a. Transparent Statistical Methods

In part 1, the researchers used a Generalized Linear Model (GLM) to reach their conclusion and compute that p-value, since the claim answers the central question of how distance to water edge predicts the probability of American Avocet microhabitat use. Additionally, we can make this conclusion because the variables are binary (use of microhabitat or no use) and continuous (distance from water edge) are non-parametric. Furthermore, since the response variable is binary we know the test was a logistic regression, a specific type of generalized linear regression. In part 2, they used a chi-square test because they are examining the relationship between two categorical variables, bird species and habitat type.

b. Suggestions for rewriting

Part 1.

We found that American Avocets do not use microhabitats in a random fashion, but rather with respect to distance to water edge. In this sense, distance to a water edge predicts the likelihood American Avocet use a microhabitat.

Part 2.

When looking at the distribution of American Avocet, Forster's Tern, and Black-necked Stilt, we found that each species was not spread randomly among habitats. Rather, each species presence differed based on habitat most likely due to assuming a certain niche in each habitat. This conclusion takes into consideration managed pond, non-tidal marsh, salt pond, tidal marsh, and tidal mudflat as habitats.

c. Further Analysis

Another variable that could influence the probability of American Avocet microhabitat use is the predator abundance in a given habitat, which could be discrete numerical count data or ordinal categorical with pre determined levels from low abundance to high abundance of predators. This variable will influence American Avocet habitat use because this species might avoid areas with high concentrations of predators as they pose a risk for injury, mortality, and reduce reproductive success.

Problem 2. Data Analysis

a. Clean and Wrangle

```
# create new object, nest_clean
nests_clean <- nests |>
  # keep only height, bird species, and ant species
  select(height_cm, case_control, ant_species) |>
  # filter to include on C. mimosae, C. sjostedti, and C. nigriceps
  filter(ant_species %in% c("AB", "RRB", "BBR")) |>
  # create new column for scientific species names
  mutate(scientific_name = case_match(
    ant_species,
    "AB" ~ "Crematogaster sjostedti", #C. sjostedti
    "RRB" ~ "Crematogaster mimosae", # C. mimosae
    "BBR" ~ "Crematogaster nigriceps" # C. nigriceps
  )) |>
  # order the species
  mutate(scientific_name =
    as_factor(scientific_name), # factor by name
    # level by scientific name
    scientific_name = fct_relevel(scientific_name,
                                   "Crematogaster sjostedti",
                                   "Crematogaster mimosae",
                                   "Crematogaster nigriceps"))

# display new structure of the data
str(nests_clean)
```

```
tibble [270 x 4] (S3: tbl_df/tbl/data.frame)
 $ height_cm      : num [1:270] 392 310 513 154 324 ...
 $ case_control   : num [1:270] 1 0 0 0 0 1 0 0 0 1 ...
 $ ant_species    : chr [1:270] "RRB" "RRB" "RRB" "RRB" ...
 $ scientific_name: Factor w/ 3 levels "Crematogaster sjostedti",...: 2 2 2 2 1 2 3 2 1 3 ...
```

```
# display 10 random rows from data
slice_sample(nests_clean, n =10)
```

```
# A tibble: 10 x 4
  height_cm case_control ant_species scientific_name
  <dbl>      <dbl>      <chr>      <fct>
1  392        1      RRB      Crematogaster sjostedti
2  310        0      RRB      Crematogaster sjostedti
3  513        0      RRB      Crematogaster sjostedti
4  154        0      RRB      Crematogaster sjostedti
5  324        1      RRB      Crematogaster sjostedti
6  392        0      RRB      Crematogaster sjostedti
7  310        0      RRB      Crematogaster sjostedti
8  513        0      RRB      Crematogaster sjostedti
9  154        0      RRB      Crematogaster sjostedti
10 324        1      RRB      Crematogaster sjostedti
```

	<dbl>	<dbl>	<chr>	<fct>
1	180.	1	RRB	Crematogaster mimosae
2	141	0	BBR	Crematogaster nigriceps
3	410	0	RRB	Crematogaster mimosae
4	123	0	RRB	Crematogaster mimosae
5	180	0	RRB	Crematogaster mimosae
6	196.	0	RRB	Crematogaster mimosae
7	270.	0	RRB	Crematogaster mimosae
8	298.	0	RRB	Crematogaster mimosae
9	158.	0	RRB	Crematogaster mimosae
10	390	1	BBR	Crematogaster nigriceps

b. Response Variable

The 1s and 0s in the case_control column indicates whether or not a tree contained a nest. If a data point has a 1 this indicates the tree contains a nest, whereas a 0 indicates the tree did not contain a nest and was “available” for comparison.

c. Purpose of Study

Tree Height

Tree height is a continuous variable measured in cm. This variable can influence the probability of nest occurrence as tree height increases the probability of nest occurrence also increases because they might provide more shelter from predators or provide necessary vantage points to identify food sources. This is supported by the fact that tree architecture plays an important role in nest site selection (Discussion, Paragraph 2) and that after running a logistic regression model it was found that each 1 m increase in height increased the odds of a bird nesting in a tree by 20% (“Results”, Paragraph 2)

Acacia Ant Species

Acacia Ant Species is a categorical variable, broken into three species *C. mimosae*, *C. sjostedti*, and *C. nigriceps*. The articles discusses how the presence of aggressive ants could cause decreasing in the number of nests due to the fact that these are nest predators and might inhibit their ability to find food (“Introduction”, Paragraph 3); however, they also discuss how the presence of aggressive ants might increase the probability of nesting occurrences because they ward off predators making these trees safer for bird to live and raise offspring (“Introduction”, Paragraph 3)

d. Table of Models

Model number	Ant Species	Tree Height	Model description
0	X	X	All predictors (No Interaction term, +)
1	X	X	All predictors (Interaction term, *)

e. Run the Models

```
# model 0: all predictors, no interaction term
nest_model0 <- glm(case_control ~ scientific_name + height_cm,
  # no interaction model variables
  data = nests_clean, # use clean dataset
  family = binomial) # set to binomial

# model 1: all predictors, interaction term
nest_model1 <- glm(case_control ~ scientific_name * height_cm,
  # interaction model variables
  data = nests_clean, # use clean data set
  family = binomial) # set to binomial
```

f. Select the Best Model


```
# calculate AIC for two models
AICc(
  nest_model0, # model 0 no interaction
  nest_model1 # model 1 interaction
) |>
  # arranging outputs in order from lowest to highest
  arrange(AICc)
```

```
      df      AICc
nest_model1  6 238.1236
nest_model0  4 258.8940
```

The best model for predicting the probability of bird nest presence as determined by Akaike's Information Criterion (AIC) is the model of interaction between tree height (cm) and ant species, demonstrated by the lower AIC value.

g. Check the Diagnostics

```
# check diagnostic with DHARMA
plot(
  simulateResiduals(nest_model1) # set to model 1
)
```



The QQ plot indicates that there is a uniform distribution of the data, since the points follow the reference line. On the right is the residuals vs. predicted plots. Since the residuals have no distinct pattern, they are randomly scattered and have a random distribution of residuals. The assumption to use a generalized linear model is met

h. Visualize the Model Predictions

```
# create preds object
preds <- ggpredict(
  nest_model1, # model object 1
  terms = c("height_cm[all]", "scientific_name")
  # predictors height and scientific name
```

```

)

# add scientific_name column for wrap in visualization
preds$scientific_name <- preds$group


# base layer: ggplot
ggplot(data = nests_clean, # set data to nest_clean
       mapping = aes(x = height_cm, # x-axis height
                     y = case_control, # y-axis nest occurrence
                     color = as.factor(scientific_name))) +
# color based on ant name
# relabelling x- and y-axis
labs(x = "Height of Acacia (cm)", # x-axis height
     y = "Probability of Bird Nest Occurrence", # y-axis nest probability
     color = "scientific_name", # color by ant species
     title =
       "Nest Occurrence Probability by Ant Species & Tree Height(cm)") +
# add title
# manually change ant colors
scale_color_manual(values = c("Crematogaster sjostedti" = "lightblue",
                              "Crematogaster mimosae" = "firebrick",
                              "Crematogaster nigriceps" = "darkgreen")) +
# fill color based on species
scale_fill_manual(values = c("Crematogaster sjostedti" = "lightblue",
                              "Crematogaster mimosae" = "firebrick",
                              "Crematogaster nigriceps" = "darkgreen")) +
# change theme from default
theme_classic() +
# remove gridlines and background
theme(panel.background = element_rect()) +
# remove legend
theme(legend.position = "none") +
# first layer: geom points, showing underlying data
geom_point(mapping = aes(x = height_cm, # x-axis: acacia height
                        y = case_control, # y-axis: nest occurrence
                        color = scientific_name), # set to selected colors
           shape = 21, # change shape
           alpha = 0.4, # make slightly transparent
           inherit.aes = FALSE) + # ignore global aes for colors
# second layer: geom line, showing predictions
geom_line(data = preds, # data from model 1 prediction
          mapping = aes(x = x, # x-axis height of acacia

```

```

        y = predicted, # y axis: prediction probability
        color = group), # color by ant
    linewidth = 1) +
# third layer: ribbon (confidence interval)
geom_ribbon(data = preds, # data from model 1 prediction
    mapping = aes(x = x, # x-axis height of acacia
        ymin = conf.low, # set low value
        ymax = conf.high, # set high value
        fill = group), # fill by ant species
    inherit.aes = FALSE, # ignore global aes
    alpha = 0.1) + # make transparent
# make new panel for each ant species
facet_wrap(~ scientific_name)

```



i. Write a Caption for your Figure

Figure 1. Probability of bird nest occurrence differs depending on the ant species present and acacia tree height Each figure separates the data, the prediction of nest occurrence given tree height (cm), and the 95% confidence interval based on ant species. The blue figure depicts data associated with the *Crematogaster sjostedti*, the red figure depicts data associated with *Crematogaster mimosae*, and the green figure represents data

associated with *Crematogaster nigriceps*. Circles represent individual observations of nest occurrences in each scenario, the lines represent model predictions of the probability of nest occurrences given tree height (cm) and ant species, and the ribbon represents the 95% confidence interval. Data originally from, Lujan, E., et al. (2023). Data, code, and metadata for: Symbiotic acacia ants drive nesting behavior by birds in an African savanna. Zenodo. <https://doi.org/10.5281/zenodo.8373322>

j. Calculate Model Predictions

```
# running ggpredict for model prediction at height of 600cm
ggpredict(nest_model1, # use model 1 w/ interaction
          terms = c("height_cm [600]", # 1st term = height 600 cm
                    "scientific_name")) # 2nd term = ant species
```

```
# Predicted probabilities of case_control
```

```
scientific_name: Crematogaster sjostedti
```

height_cm	Predicted	95% CI
600	0.20	0.00, 0.99

```
scientific_name: Crematogaster mimosae
```

height_cm	Predicted	95% CI
600	0.74	0.48, 0.90

```
scientific_name: Crematogaster nigriceps
```

height_cm	Predicted	95% CI
600	1.00	1.00, 1.00

k. Interpret your Results

According to the results from part j., at tree height of 600 cm, the probability of bird nest occurrence with *Crematogaster sjostedti* present is 0.2 (95% CI: [0.00, 0.99]), with *Crematogaster mimosae* present is 0.74 (95% CI:[0.48, 0.90]), and with *Crematogaster nigriceps* present is 1.00 (95% CI: [1.00, 1.00]). According to both figure 1. and the prediction probabilities from the

output of part j., there is a relationship between birds' preference to nest in taller trees inhabited by *Crematogaster mimosae* and *Crematogaster nigriceps* as nesting probability increases steeply with tree height for both species. Figure 1. demonstrates how the prediction probability curve increases steeply as tree height increases for *C. mimosae* and *C. nigriceps* which is also reflected in the predicted probabilities being large values of 0.74 and 1.00 respectively at 600 cm; however, when observing the *C. sjostedti* the figure remains relatively flat across all heights and only reaches a probability of 0.2 at 600 cm in part j. Biologically, this preference for trees inhabited by aggressive ant species (*Crematogaster mimosae* and *Crematogaster nigriceps*) occurs because the ants do not directly interfere with adults and offspring and actually reduce the risk of bird predation, whereas the *Crematogaster sjostedti* only provides minimal protection. There is a greater preference towards *Crematogaster nigriceps* because this species prunes, shortening shoots and creates defensive canopies that provide further protection for birds from predators like snakes, mesocarnivores, and raptors.

Problem 3. Affective and Exploratory Visualizations

a. Comparing Visualizations

How are the visualizations different from each other in the way you have represented your data?

My affective visualization was extremely different from my initial exploratory visualization in Homework 2 as the data represented and the formats were drastically altered. I originally chose to create a histogram showing count data of where dinner was cooked over the different days of the week and the background scatterplot for a generalized linear model that demonstrated the probability of cooking at home on the y axis and hours worked on the x axis; however, my affective visualization created individual figures that depicted the hours worked and whether I cooked at home or not depending on that specific day. While my original visualization had no connection between time, whether I cooked at home or ate out, and hours I worked, the affective visualization was a calendar and incorporate data on the day of the week, with images that depicted hours worked and whether I cooked or ate out.

What similarities do you see between all your visualizations?

At its core, these visualizations depict individual data points that give insight of how many hours I worked in one day and whether or not I cooked dinner at home that day. Whether its a open circle on the scatter plot background of a GLM model or the number of cracks on one plate holding dinner or a takeout bag, these styles reflect an overall trend that connects my hours working with whether I cooked at home or not.

What patterns (e.g. differences in means/counts/proportions/medians, trends through time, relationships between variables) do you see in each visualization? Are these different between visualizations? If so, why? If not, why not?

Both visualizations show the same pattern, as hours worked increases the proportion of days I ordered food out increases (and cooking at home decreases). In the explanatory visualization, this is shown by most data points on the probability line of 1.0 being clustered around low hours worked and high counts on the 0.0 line clustering at high hours worked, whereas my affective visualization portrayed a similar relationship through increased cracks on plates (hours worked) predominantly on plates with takeout bags. The patterns seem similar to one another because both draw from the same data set, the only difference is the style in which I communicate these specific data points.

What kinds of feedback did you get during week 9 workshop or from the instructors? How did you implement or try those suggestions? If you tried and kept those suggestions, explain how and why; if not, explain why not?

During workshop 9 my main feedback was to include more of the explanatory variables that I kept track of in my affective visualization, rather than just using hours worked and whether I cooked dinner or not. I implemented these suggestions by adding more novel characteristics to my calendar like color of the plate to represent whether I grocery shopped and number of prongs on a fork to indicate how many assignments I had due that day. I ended up keeping these suggestions by recreating my visualization and drawing these factors in with colored pencil because I felt as though they added more depth to my research question and demonstrated how other variables might have an affect on my dinner habits, not just hours I worked.

b. Designing an Analysis

My primary research question was “How does daily work hours relate to my decision to cook dinner at home or eat out?” My predictor variable was hours worked which is a continuous variable and my response variable was whether or not I cooked at home, an example of categorical, binary data.

Given my central question and my variables of interest, an analysis that I could apply to my research question is a generalized linear model with a binomial error distribution. This analysis makes sense because my central question is examining how one variable (hours worked) affects the likelihood of another variable (whether I cook at home or not). Additionally, since the response variable is binary a logistic regression fits best.

c. Check any Assumptions and Run your Analysis

```
# store dinner data as dinner_clean
dinner_clean <- dinner |>
  # clean names, making lowercase and underscore
  clean_names() |>
  # name treatments
  mutate(cooked_at_home_y_n = as.numeric(case_when(
    # mutate yes to "1"
    cooked_at_home_y_n == "Yes" ~ "1",
    # mutate no to "0"
    cooked_at_home_y_n == "No" ~ "0"))))
```

As a small side note, I would like to acknowledge that when running GLMs it is common practice to make model tables to demonstrate the various interactions between variables and run an AIC to find which model fits the data best; however, after speaking with Professor Bui I determined since I am only examining the affect between two variables these were unnecessary.


Additionally, since my data points were independent observations my data meets that assumption for ther analysis. Furthermore, the response variable is binomial meaning it is distributed following an exponential family distribution, meeting another assumption

Write in Model

```
# create model for where I cooked and time worked
dinner_mod <- glm(cooked_at_home_y_n ~ time_worked_hours,
  # set variables of interest
  data = dinner_clean, # set data to dinner_clean
  family = "binomial") # binomial family distribution
```

Check the Diagnostics with DHARMA

```
# run diagnostic check using DHARMA package
plot(
  simulateResiduals(dinner_mod) # use dinner model
)
```



The DHARMA diagnostic, demonstrates through the QQ plot that the residuals follow a uniform distribution since the points follow the red guideline. Additionally the dispersion test and outlier test, tell us no residuals experience a significant amount of dispersion or act as outliers impacting results. Additionally the right figure tells us that no significant problems were detected between the residuals and predicted. The assumptions to use a generalized linear model are met.

Model Coefficients and Calculating Probabilities

```
# create a summary table of models coefficients
summary(dinner_mod)
```

Call:

```
glm(formula = cooked_at_home_y_n ~ time_worked_hours, family = "binomial",
     data = dinner_clean)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	1.8462	0.7069	2.612	0.00901 **

```
time_worked_hours -0.3039      0.1479 -2.055  0.03984 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

(Dispersion parameter for binomial family taken to be 1)

```
Null deviance: 52.644  on 40  degrees of freedom
Residual deviance: 47.808  on 39  degrees of freedom
AIC: 51.808
```

Number of Fisher Scoring iterations: 4

```
# calculate probability of cooking at home when I work 0 hours
plogis(1.8462)
```

```
[1] 0.8636803
```

```
# calculate odds ratio
exp(-0.3039)
```

```
[1] 0.7379347
```

```
# confidence intervals of odds ratio
exp(confint(dinner_mod))
```

```

                2.5 %    97.5 %
(Intercept)    1.7926414 30.2930849
time_worked_hours 0.5356862 0.9687707
```

```
# calculate the probability of cooking at home after working 8 hours
ggpredict(dinner_mod, # set model
          terms = "time_worked_hours [8]") # set hours worked to 8
```


Predicted probabilities of cooked_at_home_y_n

```
time_worked_hours | Predicted |    95% CI
-----
            8 |      0.36 | 0.13, 0.67
```

d. Create a Visualization

```
# create pred_dinner object
preds_dinner <- ggpredict(
  dinner_mod, # use dinner model
  terms = c("time_worked_hours [all]")
) # predictor is time worked

# base layer: ggplot
ggplot(data = dinner_clean, # data from dinner_clean
       mapping = aes(x = time_worked_hours, # x-axis time worked
                     y = cooked_at_home_y_n)) + # y-axis whether I cooked
# relabelling x- and y- axes
labs(x = "Time Worked (Hours)", # x-axis time worked
     y = "Probability of Cooking at Home", # y-axis probability of cooking
     title = "Relationship between Hours Worked and my Cooking Habits") +
# create title
# first layer: geom points showing data
geom_point(color = "steelblue", # change color
           alpha = 0.4, # change opacity
           shape = 21) + # make open circles
# second layer: geom line showing probability
geom_line(data = preds_dinner, # set data to prediction
          mapping = aes(x = x, # x variable
                        y = predicted), # y variable is predicted
          color = "firebrick", # change color
          linewidth = 1, # alter line width
          inherit.aes = FALSE) + # don't use ggplot aes
# third layer: ribbon (95% CI Prediction)
geom_ribbon(data = preds_dinner, # data from prediction
           mapping = aes(x = x, # set x axis data
                         ymin = conf.low,
                         # lower confidence interval for ribbon
                         ymax = conf.high),
           # upper confidence interval for ribbon
           fill = "firebrick", # color of CI
           alpha = 0.4, # opacity of CI
           inherit.aes = FALSE) + # don't use ggplot aes
theme_minimal() # change theme
```



e. Write a Caption

Figure 2. Probability of cooking dinner differs based on hours worked. The figure depicts the individual data points, predicted probability based on hours worked, and the confidence interval associated with the probability that I cook at home depending on how much I work. 1 indicates I cook at home and 0 indicates I eat out. Small blue open circles represent individual observations of whether I cooked at home or not in relation to the amount I worked that day. The bold red line represents the model prediction of cooking at home given time worked(hours). Finally the red ribbon represents the 95% confidence interval. Data collected by Stephen Bond from 2026-04-20 to 2026-05-30.

f. Write About Your Results

The probability of cooking at home when time worked is 0 hours is 0.86. With every 1 hour increase in time worked, the odds of me cooking at home decrease by a factor of 0.74(95% CI:[0.536, 0.969], $p = 0.04$, $\alpha = 0.05$). When working 8 hours, the probability of me cooking at home is 0.36(95% CI:[0.13, 0.67]). The model coefficient of 0.74 from the odds ratio summary along with the predicted probability curve from Figure 2. decreasing from 0.86 at 0 hours to 0.36 at 8 hours, indicate that I tend to cook dinner at home more often when I work fewer hours.

g. Sharing Your Affective Visualization

I was present in workshop!